

G1 Recovery AMP Get-Up — 8192-Env Fresh 10k

Training report for RECOVERY_AMP_GETUP_8192ENV_FRESH_10K

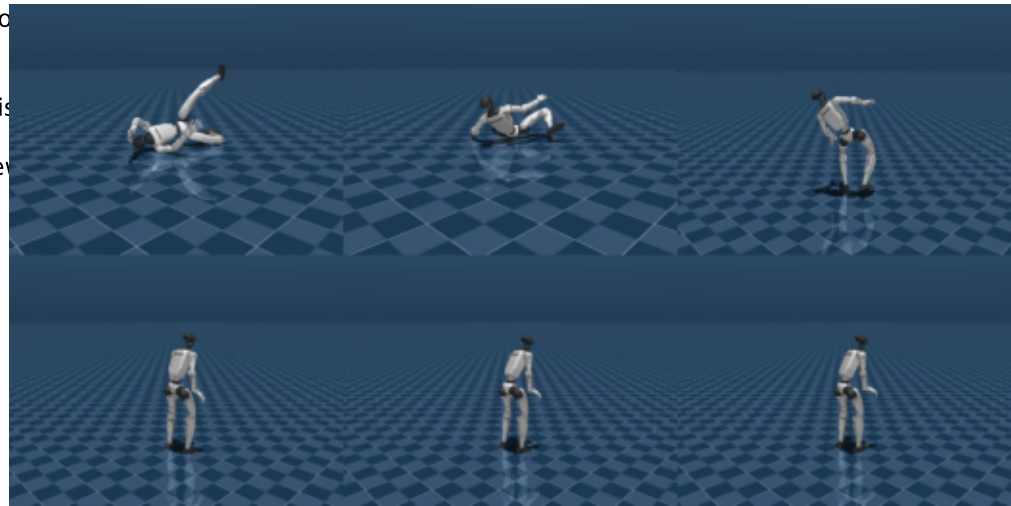
Logged through iteration 9,999 · approximately 7.6 h wall-clock · 2 GPUs · 4,096 environments/rank

Executive assessment

- The requested run reached its 10k endpoint and produced model 9999.pt.

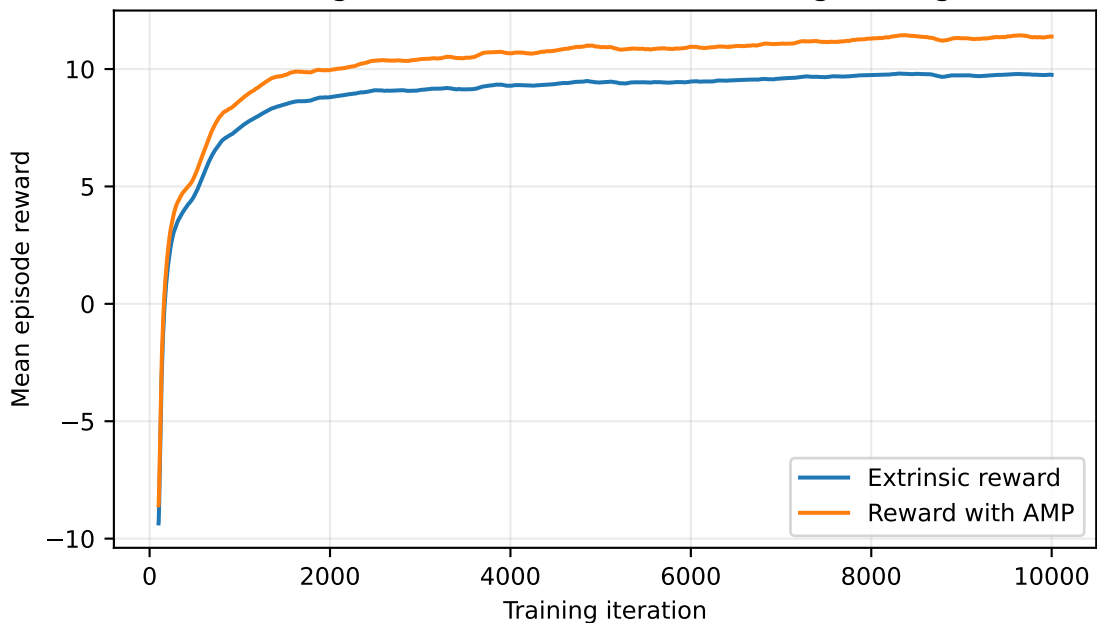
Final model_9999 MuJoCo playback — one frame per second

- The final MuJoCo side-fall video shows a successful transition through the clip.
- This is evidence for the deterministic side-fall case only. It is
- Training metrics should be judged by trend and stability; rev

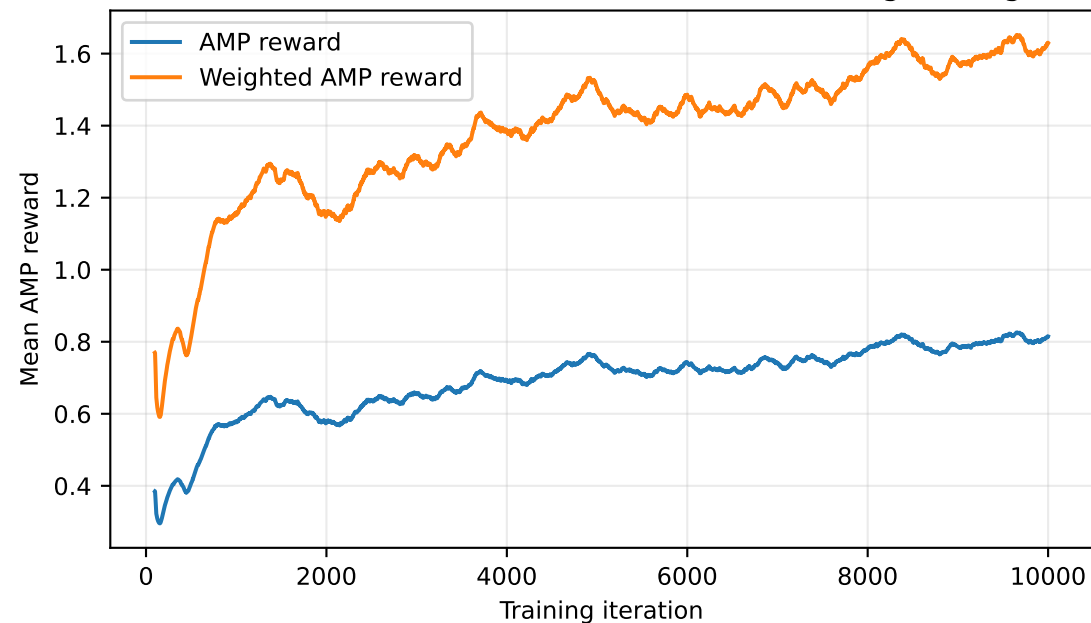


Learning progress

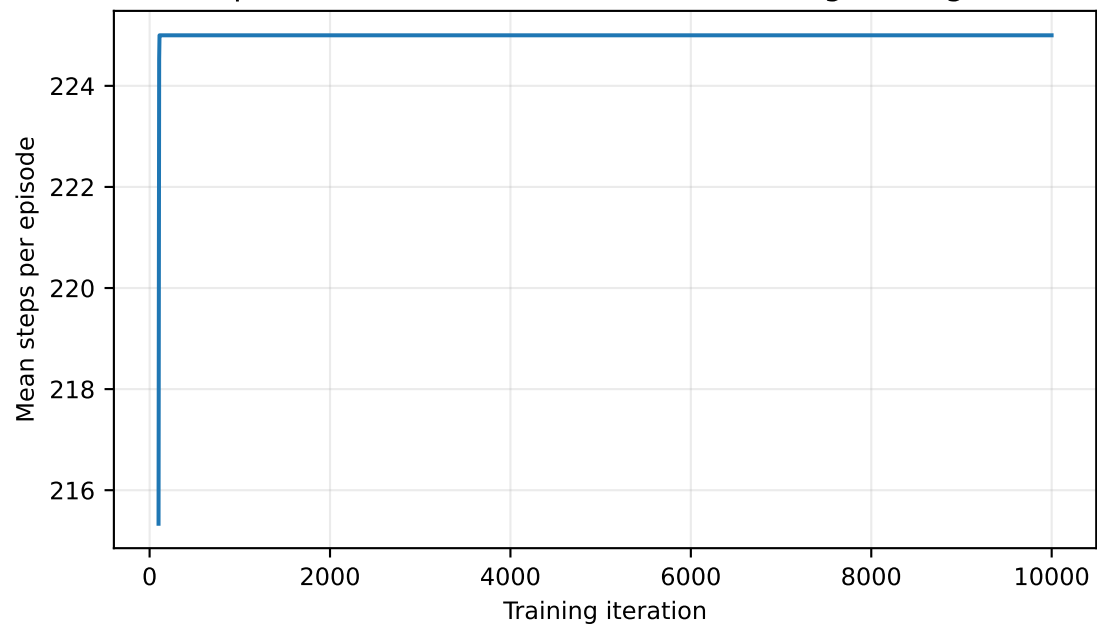
Training reward (100-iteration moving average)



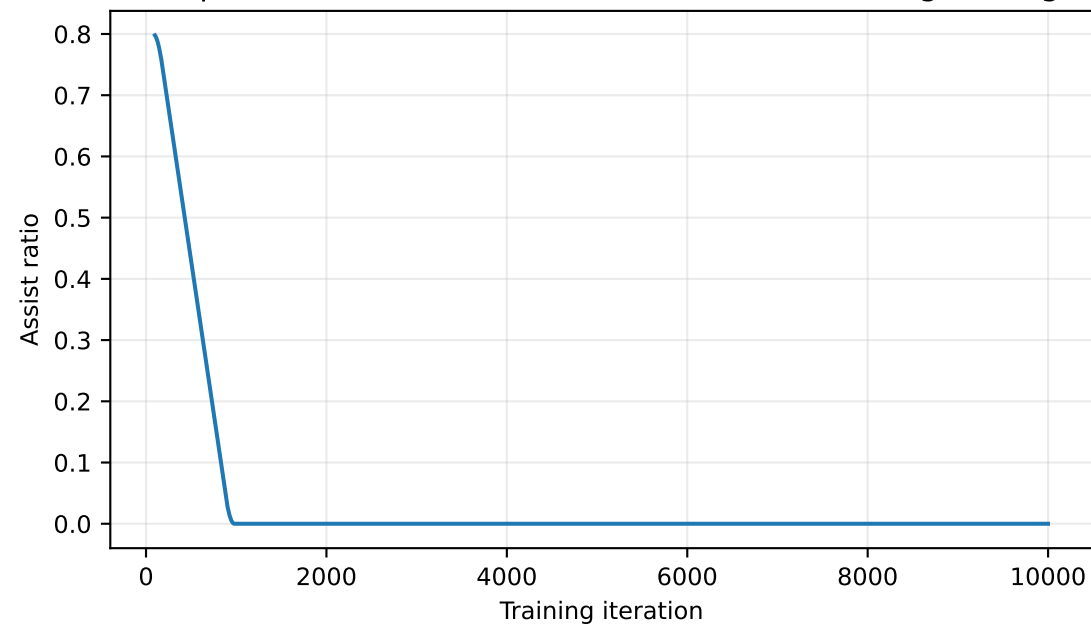
AMP reward contribution (100-iteration moving average)



Episode duration (100-iteration moving average)

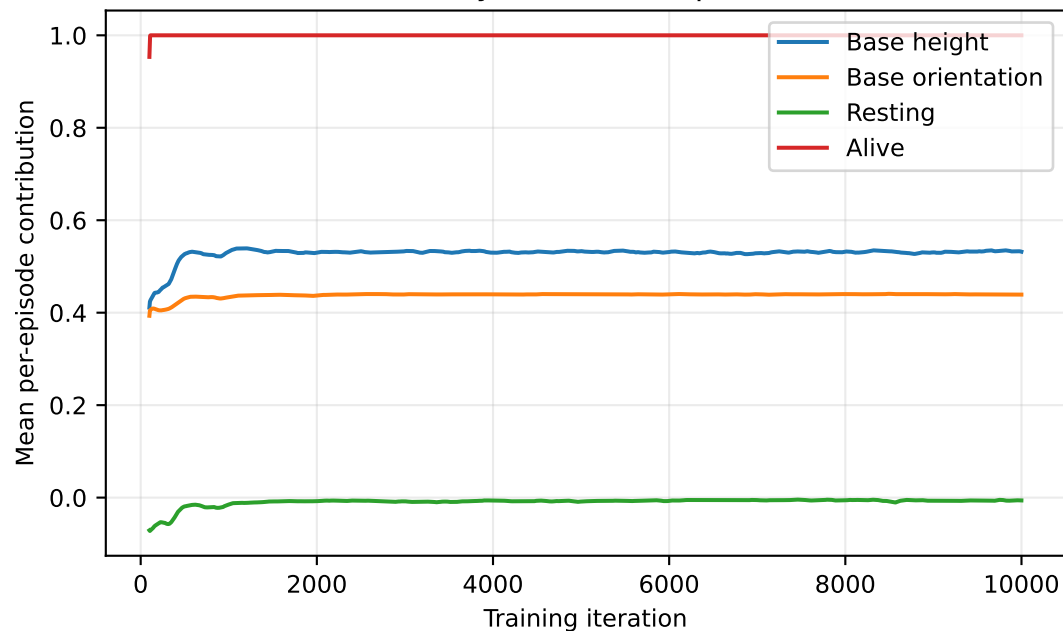


Get-up assistance curriculum (100-iteration moving average)

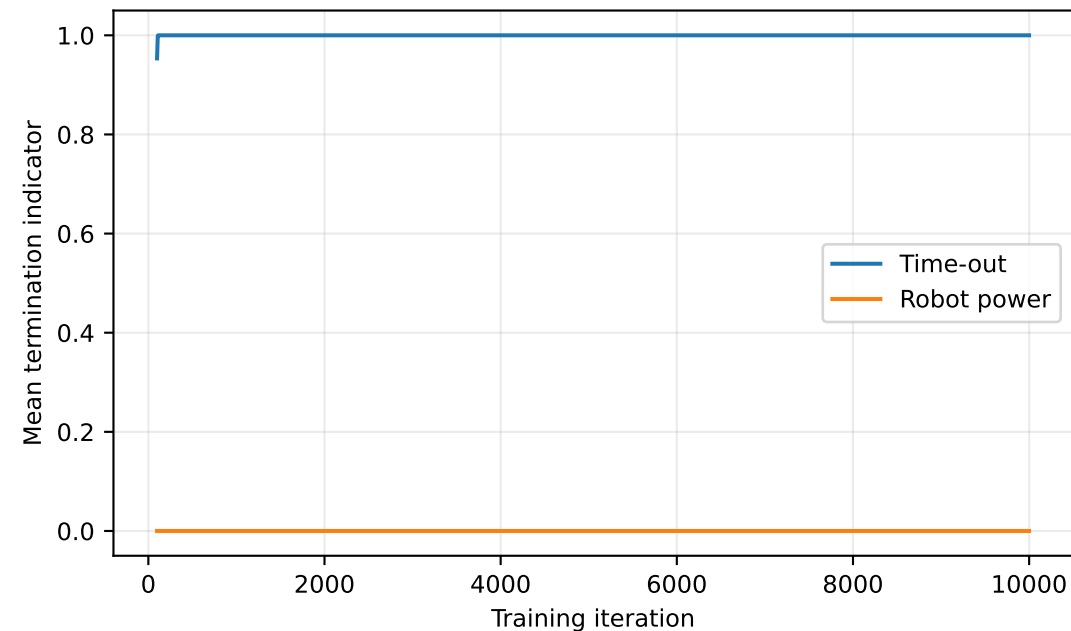


Recovery behavior and efficiency

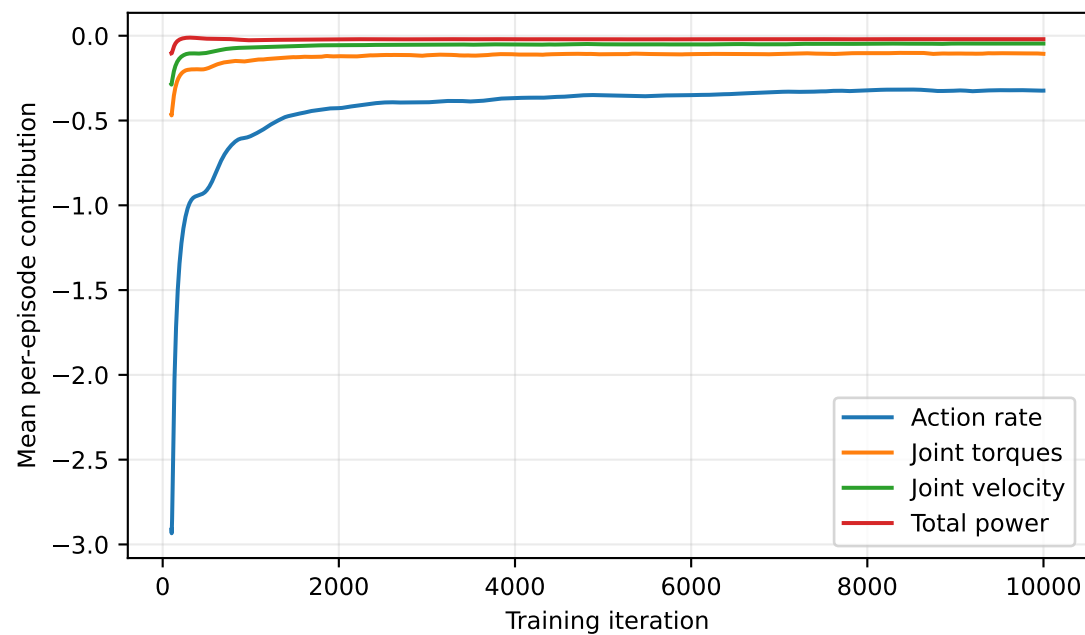
Recovery reward components



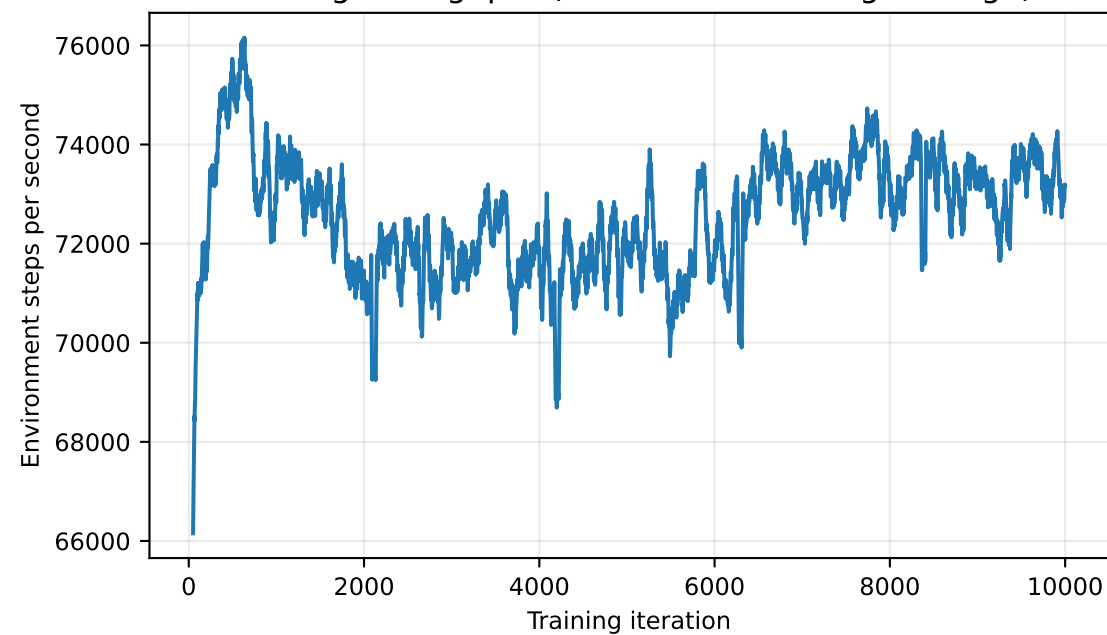
Termination rates



Control-cost reward terms

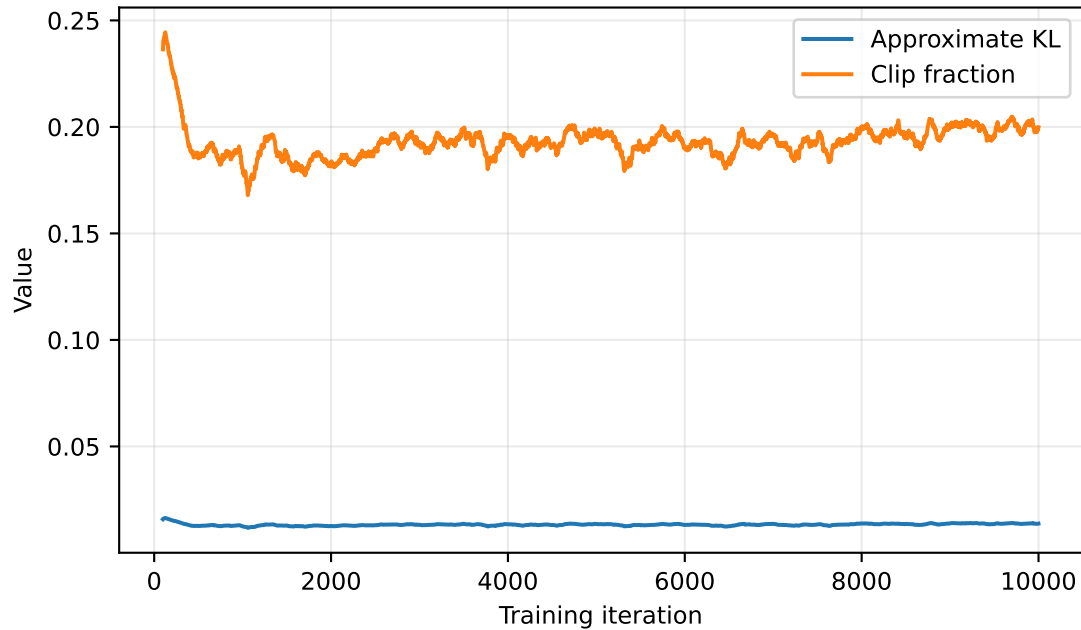


Training throughput (50-iteration moving average)

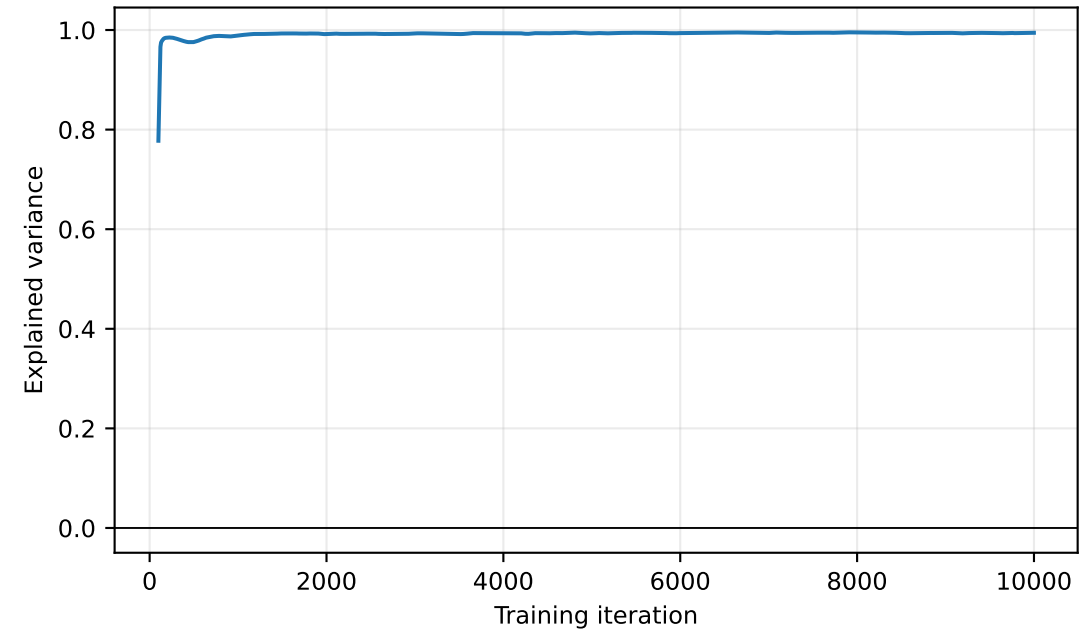


Optimization health

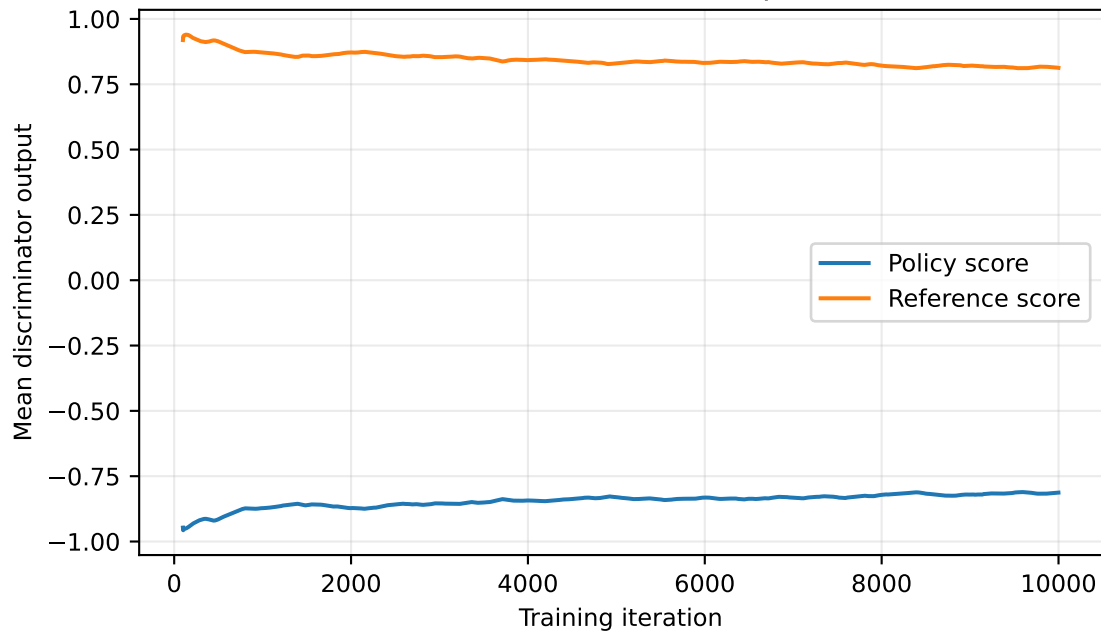
PPO update size



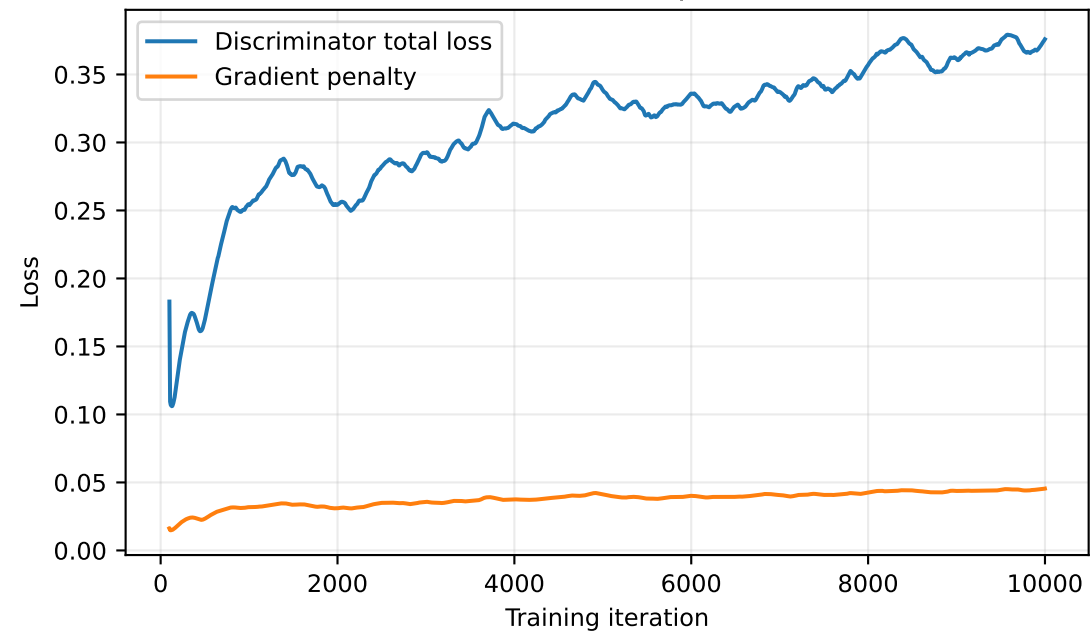
Critic explained variance



AMP discriminator outputs



AMP discriminator optimization



Selected metric summary

Train/mean_reward

first: -2.638
final: 9.657
last-100 mean: 9.753
best: 9.951 @ 8348

Train/mean_reward_with_amp

first: -2.023
final: 11.200
last-100 mean: 11.384
best: 11.775 @ 9916

Train/mean_amp_reward

first: 0.307
final: 0.772
last-100 mean: 0.815
best: 1.047 @ 8

Train/mean_episode_length

first: 21.110
final: 225.000
last-100 mean: 225.000
best: 225.000 @ 9

Curriculum/get_up_assist_ratio/get_up_assist_ratio

first: 0.800
final: 0.000
last-100 mean: 0.000
best: 0.000 @ 880

Perf/total_fps

first: 44403.000
final: 74096.000
last-100 mean: 73110.710
best: 80283.000 @ 494

Episode_Reward/recovery_base_height

first: 0.059
final: 0.529
last-100 mean: 0.532
best: 0.547 @ 1167

Episode_Reward/recovery_base_orientation

first: 0.035
final: 0.439
last-100 mean: 0.439
best: 0.442 @ 2443

Episode_Reward/resting

first: -0.008
final: -0.006
last-100 mean: -0.006
best: -0.002 @ 7740

Loss/approx_kl

first: 0.005
final: 0.013
last-100 mean: 0.014
best: 0.005 @ 0

Loss/clip_fraction

first: 0.052
final: 0.191
last-100 mean: 0.200
best: 0.052 @ 0

Loss/explained_variance

first: 0.033
final: 0.995
last-100 mean: 0.994
best: 0.996 @ 7763

Interpretation limits: TensorBoard rewards are task-specific training signals. The included video is one deterministic MuJoCo side-fall rollout and does not establish a randomized recovery success rate, hardware robustness, or sim-to-real readiness.